

# CASI day 5 lecture notes: Neural networks

*Tuomo Nieminen*

*May 16, 2018*

## Neural networks

### Introduction

The goal of machine learning is to find a function  $f$  to characterize a predictive relationship between the data  $X$  and some target  $y$ , i.e. to find an  $f$  for which

$$f(X) = y.$$

A simple example is linear regression, where

$$f(X) = w_0 + \sum_{j=1, \dots, p} x_j w_j.$$

### Neural networks

A neural network is a highly parametrized model which in theory can learn any smooth predictive relationship  $f : X \rightarrow y$ .

### Layers

A neural network consists of  $K$  layers ( $L_k$ ), which can be categorized to

- An input layer ( $L_1$ )
- Hidden layers ( $L_2, \dots, L_{K-1}$ )
- An output layer ( $L_K$ )

*See fig. 18.1. in Computer age statistical inference*

### Units and activations

The layers of a neural network consist of *units*, which have *activation functions*. The outputs of the activation functions are fed to the next layer. The unit  $l$  in layer  $k$  takes as input the outputs from the previous layer, combines them via a linear function, and then applies a nonlinear transformation  $g_l$  to produce outputs  $a_l^{(k)}$ .

So, for example the output of the units in the second layer are nonlinear transformations of linear combinations of the input vector:

$$z_l^{(2)} = w_{l0}^{(1)} + \sum_{j=1, \dots, p} w_{lj}^{(1)} x_j$$

$$a_l^{(2)} = g_l^{(2)}(z_l^{(2)}).$$

## Transitions between layers

In general, the transition  $L_{k-1} \rightarrow L_k$  is given by

$$\begin{aligned}z^{(k)} &= W^{(k-1)} a^{(k-1)} \\ a^{(k)} &= g^{(k)}(z^{(k)})\end{aligned}$$

where  $W$  represents the matrix of weights from layer  $L_{k-1}$  to layer  $L_k$ ,  $a^{(k)}$  is the entire vector of activations at layer  $L_k$ , and  $g^{(k)}$  operates elementwise.

*ref: Computer age statistical inference, pages 351 - 356.*

## Activation functions

Typical choices for the nonlinear part of the activation functions are

- Rectified linear unit (ReLU):  $g(x) = \max(0, x)$
- Sigmoid:  $g(x) = \frac{e^x}{1+e^x}$

Without the nonlinearities, the neural network reduces to a generalized linear model.

## A neural network for classification

Say we want to predict one of 10 classes, using the data  $x$  as input, i.e. we want

$$f(x) = \{0, 1, \dots, 9\}$$

It is generally more useful to learn the probability function

$$P(y = j|x), \quad j = 0, 1, \dots, 9$$

We will define a neural network which produces probabilities for each of the labels.

## Output layer

In this example the output layer of the network has 10 nodes, one for each of the possible class labels. The final nonlinear transformation  $g^{(K)}$  is the *softmax function*

$$g^{(K)}(z_m^{(K)}; z^{(K)}) = \frac{\exp(z_m^{(K)})}{\sum_{l=1, \dots, M} \exp(z_l^{(K)})},$$

which computes a number (probability) between zero and one, and all  $M$  of them sum to one. This is a symmetric version of the inverse link function used for multiclass logistic regression: a generalization of the sigmoid.

*ref: Computer age statistical inference, page 355*

## Fitting a neural network

### Train and test data

A machine learning algorithm is usually trained by first splitting the data into train and test sets. The training data is used to fit the model and the test data is used to evaluate the model performance.

### Loss function

A neural network model is a hierarchical function  $f(x; W)$  of the feature vector  $x$ , and the collection of weights  $W$ . A loss function is a function defined on a single data point  $x_i$ , a prediction  $\hat{y}_i = f(x_i|W)$  and target  $y_i$ , and measures the penalty caused by an error in the prediction. For example the square loss is given by

$$L(f(x_i|W), y_i) = (f(x_i|W) - y_i)^2 = (\hat{y}_i - y_i)^2$$

### Cost function

A cost function defines the optimization problem to determine the optimal set of parameters in the model. A typical choice is the average loss over the training set plus some model complexity penalty (regularization):

$$C(W) = \frac{1}{n} \sum_{i=1, \dots, n} L(f(x_i|W), y_i) + \lambda J(W),$$

where  $\lambda$  is a tunign parameter for the penalty function  $J(W)$ . For example, using the square loss and no penalty ( $\lambda = 0$ ), the cost becomes

$$\frac{1}{n} \sum_{i=1, \dots, n} (\hat{y}_i - y_i)^2,$$

also known as Ordinary least squares (OLS).

### Objective function

A cost function can be derived via maximum likelihood estimation (possibly adding a penalty term). An MLE is a type of objective function, a general expression for the function to optimize.

ref: [stats.stackexchange.com](https://stats.stackexchange.com)

### Gradient Descent

To fit the model, the objective is to minimize the cost function with respect to  $W$ . There is no analytical solution, so we resort to numerical methods. Given the gradient  $\Delta C(W)$  of the cost function, the gradient descent update on the weights is given by

$$W^{(k)} \leftarrow W^{(k)} - \alpha \Delta C(W^{(k)}), \quad k = 1, \dots, K,$$

where  $\alpha$  is the learning rate. Gradient descent requires (nonzero, typically random) starting values for all the weights  $W$ .

## Backpropagation to find the gradient of the cost function

Since  $f(x|W)$  is defined hierarchically, the elements of  $W$  occur in layers, starting from the input layer. Computing the gradient is also done most naturally in layers. The backpropagation algorithm implements the chain rule for differentiation to find the gradient of the objective function.

The goal is to compute the derivative of  $L(f(x|W), y)$  with respect to any of the elements of  $W$ , for a generic input-output pair  $(x, y)$ . Since the objective function is a sum of loss functions, the overall gradient will be the sum of these individual gradient elements over the training pairs  $(x_i, y_i)$

$$\Delta C(W) = \sum_{i=1, \dots, n} \Delta L(f(x_i|W), y_i) + \Delta \lambda J(W)$$

*ref: Computer age statistical inference, page 356.*

## Stochastic gradient descent

Rather than process all the observations before making a gradient step, it can be more efficient to process smaller *batches* of data at a time, or even a single data point  $i$ . The gradient computed for a batch approximates the true gradient.

$$E[\Delta C_i(W)] = \Delta C(W)$$

Using sub samples to compute the gradient is called stochastic gradient descent. Several passes can be made over the training set until the algorithm converges. Data is usually shuffled in between to avoid sticking to poor local minima (hence the term *stochastic?*). An *epoch* of training means that all  $n$  training samples have been used in gradient steps.

## Types of networks

Convolutional neural networks <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Recurrent neural networks